Feature-boundary audit pilot

Reviewer guide | FB-PILOT-v1

What you are assessing

Apply a documentary audit rubric to four assigned capabilities. Record what each source supports, which responsibilities belong to the capability, and where the rubric cannot be applied consistently. You are not testing products or reviewing the paper for acceptance.

This is an exploratory application pilot. The coordinator must approve the rules and give both reviewers the same fixed source bundle before collection starts. A URL list alone does not fix the source version. Enter the supplied corpus ID in Profile.

Independent first pass

Work in your own copy of rating_template.xlsx. Do not consult the other reviewer, the manuscript application tables, author ledgers, reviews, or a proposed answer key before returning your initial file. Record any prior exposure honestly. An author who already knows the design is not a blinded external reviewer.

Do not ask a generative AI system to select evidence or assign ratings. Record any permitted translation or other assistance. Use only the agreed corpus; log inaccessible sources and request a shared clarification rather than silently adding different material.

Order of work

1. Profile: record role, experience, prior exposure, materials and time. Use a pseudonymous reviewer ID; do not enter private company information.

2. Boundaries: interpret the assigned capability, identify its actors, configuration and responsibilities, and note scope ambiguities.

3. Evidence and Paths: give each source-to-subclaim mapping an evidence ID. Name separately evaluated paths and their effects on broader claims.

4. Ratings: complete the 20 items, linking to your evidence and path IDs. Codebook contains the anchors. Field checks indicate missing entries, not correctness.

5. Feedback: record problems with the rubric. Return an initial file before discussing disagreements. Do not overwrite it with a consensus version.

Participation

Participation and use of responses must be agreed with the coordinator in advance. Declining must not affect your employment or relationship with the authors. The coordinator is responsible for applicable institutional/company review and for consent, access and retention arrangements. This guide does not constitute ethics approval.

Assigned capability boundaries

ID and capability	Starting scope to review
B01 Apple Foundation Models	WWDC25 on-device, text-only app request/session API. Do not substitute a WWDC26 developer-selected PCC model. Describe the host-app responsibilities that lie outside this component.
B02 Apple PCC one-shot request	A single PCC request beginning after client-side request construction. Agent-specific tool paths and multi-turn processing are not part of this assigned request. Identify dependencies and separate client responsibilities.
B03 Android AICore / Gemini Nano	An app request to the local AICore inference service. App-level retrieval, actions and retained app state are not automatically part of the service. Distinguish model distribution/update paths from request processing.
B04 Microsoft Recall	Local snapshot capture, indexing and search, including the controls governing that state. Identify additional network, action and management paths and justify whether each is internal, adjacent or absent.

Configuration and source versions

Record the version, region, account, policy profile and feature flags that can actually be established from the supplied material. Write UNDETERMINED where the material does not fix an item. A source access date is not an observed product version.

B01 is the historical WWDC25 on-device API. A later framework capability is not evidence that this historical request used that route. B02 is one-shot processing, so agent-specific tool or multi-turn properties require their own scope. Do not credit or penalize one component for an adjacent component without an explicit supported dependency.

Actors and access

A0: content author. A1: host or third-party app with ordinary permissions. A2: tool recipient or external service. A3: local co-user. A4: administrator, repair or forensic actor. A5: fallback or telemetry operator. A6: local non-root malware under a specified privilege class.

Specify which classes your documentary interpretation concerns. Do not claim actor-conditioned testing. For local users or administrators, distinguish lock state, account, credentials and management authority when the sources permit it.

Source packet

Sources lists 23 candidate public-document entries and starting locators. Historical access dates are provided for orientation, not as a frozen pilot corpus. The full PCC SOC 3 report is excluded from the pilot packet because it carries a use-and-distribution notice; the public Apple overview is included. A summary saying that an examination exists is not the examination artifact itself.

The coordinator supplies the final included versions and corpus ID. Source unavailability is a coverage limitation. Do not infer a missing control from a missing document. This prospective packet is not an exact reproduction of the original manuscript corpus.

Pilot decision rules

Rating state	Use
SCORED	A supported numeric anchor can be assigned for the declared scope. Enter an integer 0-3.
UD	The dimension applies but its support or scope cannot be determined. Leave numeric score blank and explain the gap.
ADJ	The assessed path is reachable but assigned to a separately named boundary. Register it in Paths; leave score blank.
N/A	Evidence establishes that the relevant path is absent in the specified configuration. Mere exclusion or silence is insufficient. Leave score blank.

Applying anchors

Use Codebook anchor IDs, such as D1-L2, in the rationale. Wording joined by "and" requires all listed conditions; an explicit "or" permits the stated alternative. For an unresolved comma-list or mixed-ownership ambiguity, record the competing interpretations. Do not silently invent an extra requirement or force a favorable score.

Assign the highest fully supported applicable anchor, not the nearest-sounding level. Partial level-2 support does not automatically imply level 1: a level-1 condition needs its own basis. Zero requires affirmative evidence of absent or contradicted control. Missing information alone is UD.

Separate client source selection from server payload handling under D1, and observability, recourse and change obligations under D5. Record each material subclaim and owner in Evidence. If a dimension mixes internal and adjacent obligations and no unambiguous summary is supported, use UD and explain the split. ADJ for one path does not erase other internal obligations.

Do not average numeric scores across dimensions or vendors. A score records an anchor match under the stated scope; it is not measured privacy effectiveness.

Evidence grades belong to individual artifacts

E1: assertion or general commitment. **E2:** design, static code/configuration or schema. **E3:** actual execution record or completed control evaluation. **E4:** claim-matched validation package with coverage, repetition, uncertainty and change replay.

Record the issuer, artifact type, source location, covered subclaim, applicable configuration and limit. A public report can support its actual evaluated scope. A specification is not a captured log; a summary is not the underlying evaluation. Do not take the maximum grade across unrelated claims. A strong record can document a failure.

Recording and returning judgments

Evidence records

Use a separate row when the source supports a different subclaim, owner, boundary or evidence grade. Link each row to a Ratings item ID. Use the source ID from Sources, an exact heading/page/paragraph locator, and a short factual paraphrase. Do not paste long copyrighted passages or restricted material.

State whether the evidence supports an obligation inside the assigned boundary. Explain version or ownership mismatches. An inaccessible source can be recorded as a searched source with no grade; do not assign E1 just because nothing was found. More rows may be added using unique IDs.

Adjacent paths

Record the path and trigger, crossed data or authority, owner, supporting evidence, target boundary, reachability and the broader claim that remains blocked. A conditional path requires a recorded configuration condition. A target being named does not mean it has been evaluated.

List paths you identify yourself; no author-completed path register is supplied. Different wording does not necessarily mean a different path. After independent submission, the coordinator will align equivalent paths before comparing judgments.

Completing each rating

Enter a state; enter a number only for SCORED. List evidence IDs or write NONE with the sources searched recorded in Evidence. State the supported anchor and its basis, unmet or out-of-scope conditions, and any ambiguity. Write "None identified" when appropriate instead of leaving required rationale fields empty.

Ratings contains a field-completeness check. It does not verify source truth, applicability or grade. Use Feedback to identify ambiguous rules and missing concepts. If you cannot make a judgment, UD with a reason is a valid response, not a failed participant task.

No implementation outcomes in this pilot

Do not turn hypothetical examples into collected measurements. This activity does not execute products; product test outcomes remain untested. If implementation work is later conducted, privacy violations must be counted across all attempted runs even when the intended task fails. Task utility is reported separately; deny-all is not useful protection.

Return and reconciliation

Return the initial workbook under your assigned ID, for example R1_initial.xlsx. The coordinator preserves it unchanged. Discussion starts only after both initial workbooks are received. Consensus and codebook changes belong in a separate reconciliation record.

One author and one non-author colleague can provide an exploratory two-reviewer application. Their agreement does not establish independent external validation or correctness. Report relationship, experience, prior exposure, missing data, source changes and unresolved disagreements.